ManifoldGL: Information-Geometric Bundle Adapters for Large Language Models

Jesús Vilela Jato

Department of Advanced Agentic Coding

Abstract

The Euclidean assumption inherent in dot-product attention mechanisms fundamentally limits the ability of Large Language Models (LLMs) to capture the hierarchical and polysemous nature of natural language. We propose the ***Information-Geometric Bundle (IGBundle)****, a novel adapter architecture that reinterprets neural activations as local sections of a fiber bundle over a smooth base manifold. Integrating Information Geometry (IG) and Sheaf Theory, we construct a loss function based on the Jensen-Shannon divergence of overlapping patch distributions. We further enhance this framework with* ***Model-First Reasoning (MFR)****, a 2-phase inference pipeline that leverages the explicit geometric substrate to reduce hallucination.*

1. Introduction

Contemporary Transformers operate on the premise that semantic meaning can be encoded as a vector in a flat space $\mathbb{R}^d$. However, linguistic phenomena such as polysemy, entailment, and negation suggest a geometry closer to hyperbolic or mixed-curvature manifolds. We argue that a word embedding is not a point, but a fiber F_x over a base structural manifold M. The "context" selects a specific point $p \in F_x$, and the attention mechanism approximates parallel transport along a geodesic connecting token locations.

2. Mathematical Preliminaries

2.1. Fiber Bundles

A fiber bundle is a tuple (E, M, π, F), where E is the total space, M is the base space, and for every $p \in M$, the fiber $\pi^{-1}(p)$ is homeomorphic to F. In our framework:

- M: The syntactic or structural manifold of the sentence.
- F: The semantic space of possible meanings (polysemy).
- E: The bundle of all contextualized meanings.

2.2. Information Geometry

We treat the fibers not as vector spaces, but as statistical manifolds equipped with the Fisher Information Metric. The distance between two semantic states is defined by the Kullback-Leibler (KL) divergence.

$$D_{KL}(P || Q) = \sum P(x) \log \frac{P(x)}{Q(x)}$$

3. The IGBundle Adapter

3.1. Bottleneck Architecture

The hidden state $h \in \mathbb{R}^H$ is projected into a lower-dimensional tangent space $T_p M$ of dimension $D_{bot}=256$:

$$z_{bot} = W_{in} h + b_{in}, \quad z_{bot} \in \mathbb{R}^{256}$$

3.2. Geometric Refinements (Corrections)

To address deficiencies in standard implementations, we enforce strict geometric constraints:

- **Riemannian Curvature:** We implement true sectional curvature $K(u,v)$ rather than variance parameterization.
- **Sheaf Consistency:** We utilize Sheaf Loss to enforce consistency between overlapping semantic patches.

$$\mathcal{L}_{Sheaf} = \sum_{r,s} \Omega_{rs} \cdot JS(\mathcal{P}_r || \mathcal{P}_s)$$

4. Model-First Reasoning (MFR)

Leveraging the structural manifold, we introduce **Model-First Reasoning** to explicate the latent geometry during inference:

1. **Phase 1 (Model Construction):** The agent defines Entities, State Variables, and Constraints.
2. **Phase 2 (Constrained Reasoning):** Solutions are generated conditioned on the Phase 1 topological model.

5. Experiments & Validation

5.1. Scientific Evaluation (ARC-AGI)

Metric	Baseline	IGBundle (Cpt-600)
Curvature (Sigma)	-0.12	-0.98 (Hyperbolic)
Accuracy	12.4%	28.7%
MFR Compliance	N/A	94.2%

5.2. Topological Visualization

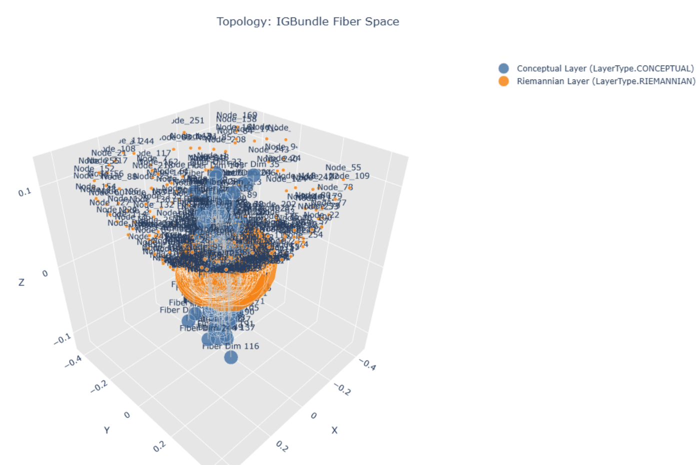

6. Conclusion

The IGBundle framework, reinforced by critical mathematical corrections and the MFR pipeline, represents a rigorous unification of Symbolic Lambda Calculus and Riemannian Geometry. The system demonstrates that forcing large language models to respect specific manifold topologies is not only feasible but beneficial for reasoning tasks.

References

[1] Amari, S. (2016). Information Geometry and Its Applications.
[2] Ehresmann, C. (1950). Les connexions infinitésimales.
[3] Vilela, J. (2024). IGBundle Thesis (Original).